

# Comparing the performance of state-of-the-art software switches in the context of NFV

Tianzhu Zhang

# About me

- 2020, 8 – now      Research Engineer @ Nokia Bell Labs
- 2017,10 – 2019,11      PostDoc @ Telecom ParisTech
- 2017, 6 – 2017, 9      Intern @ Nokia Bell Labs

# About the project

- Joint work of TPT, NBLF, and PoliTO
- Collaborators:
  - Leonardo Linguaglossa (Telecom Paris, Paris)
  - Massimo Gallo (Nokia Bell Labs, Paris-Saclay)
  - James Roberts (Telecom Paris, Paris)
  - Luigi Iannone (Telecom Paris, Paris)
  - Paolo Giaccone (Politecnico di Torino, Turin)
- Published results:
  - Demo at IEEE Netsoft'19
  - Paper at ACM CoNEXT'19
  - On-going extension

# Outline

## Background

- Network Function Virtualization (NFV)
- Software switches
- Motivation

## Methodology

- Test scenarios
- Testbed settings

## Experimental results

## Conclusion

# Background

# Ossification of traditional networks



Specialized hardware appliances, aka. Middleboxes



Very expensive



Difficult to manage



Inflexible to customize



Tightly coupled with hardware production cycle



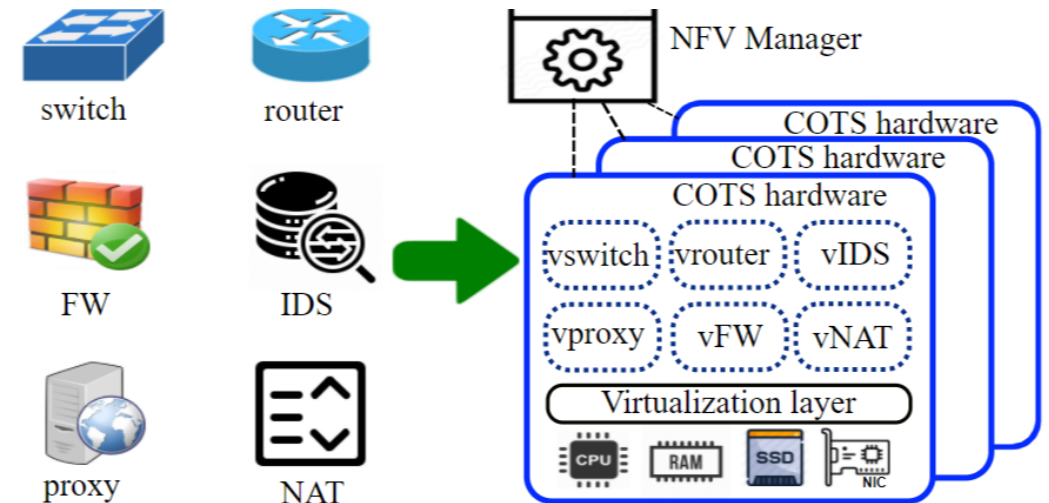
Tedious service deployment and maintenance



Non-trivial to scale in/out


# Network Function Virtualization (NFV)

- Shift towards agile and open service provisioning
  - Virtual network functions (VNFs) on Commodity Off-The-Shelf (COTS) hardware



# Software switches

- Virtual switch implemented in software
- High flexibility
- Low performance
- Previously only used for fast-prototyping and function testing
- Benefit from high-speed packet I/O techniques
  - DPDK, netmap, PF\_ring, eBPF, VMA, etc.
- Widely adopted by existing NFV frameworks
  - Traffic steering between VNFs and network interface cards (NICs)





	Architecture		Programming Paradigm	Model		Virtual Interface	Runtime Reprogrammability	Programming Language	Main Purpose
	Self-contained	Modular		RTC	Pipeline				
BESS		✓	Structured	✓	✓	vhost-user	High	C, Python	Programmable NIC
Snabb		✓	Structured		✓	vhost-user	High	Lua, C	VM-to-VM
OVS-DPDK	✓		Match/action	✓		vhost-user	High	C	SDN switch
FastClick		✓	Structured	✓		vhost-user	Medium	C++	Modular router
VPP	✓		Structured	✓		vhost-user	Medium	C	Full router
VALE	✓		Structured	✓		ptnet	Low	C	Virtual L2 Ethernet
t4p4s	✓		Match/action	✓		vhost-user	Low	C, Python	P4 switch

# Unclear performance

- Disparate design choices and purposes
- Performance depend on traffic patterns
- Too many hardware and software parameters to tune

Ref.	Software switches under test	Bare-metal	VNF environment VM      Container	Inter-VNF forwarding	SFC	Throughput Uni.    Bi.	Latency
[34]	<b>BESS</b>		✓		✓	✓	✓
[71]	<b>OVS, OVS-DPDK</b>	✓	✓      ✓ <sup>(*)</sup>		✓	✓	✓
[72]	Linux bridge, <b>VALE</b>		✓	✓	✓	✓	
[73]	ClickOS, <b>BESS</b>				✓	✓	✓
[19]	Linux bridge, OVS, <b>OVS-DPDK</b>	✓	✓	✓		✓	✓
[25]	OVS, PISCES, t4p4s	✓				✓	
[74, 75]	<b>OVS-DPDK, VPP</b>		✓			✓	
[18]	<b>BESS, VPP, OVS-DPDK</b>	✓				✓	
[76]	<b>OVS-DPDK, Snabb, VALE</b>		✓	✓		✓	
[21]	<b>Snabb, OVS, OVS-DPDK, Linux bridge</b>		✓	✓		✓	
[77, 30]	<b>OVS-DPDK, OVS, Linux bridge, Lagopus, xDPd-DPDK</b>	✓	✓			✓	✓

# Existing works


# Our contribution

- Experimental methodology to thoroughly evaluate the performance of software switches for NFV
- Compare the performance of 7 state-of-the-art high-speed software switches, quantitative-> qualitative
- Define the best use case for each tested software switch

# Methodology





# Typical NFV scenario

- Service function chain (SFC)



# Test scenarios

- Physical-to-physical (p2p)
- Physical-to-virtual (p2v)
- Virtual-to-virtual (v2v)
- Loopback



# Software switches under test

OVS-DPDK

FastClick

Berkley  
Extensible  
Software Switch  
(BESS)

Vector Packet  
Processing (VPP)

t4p4s

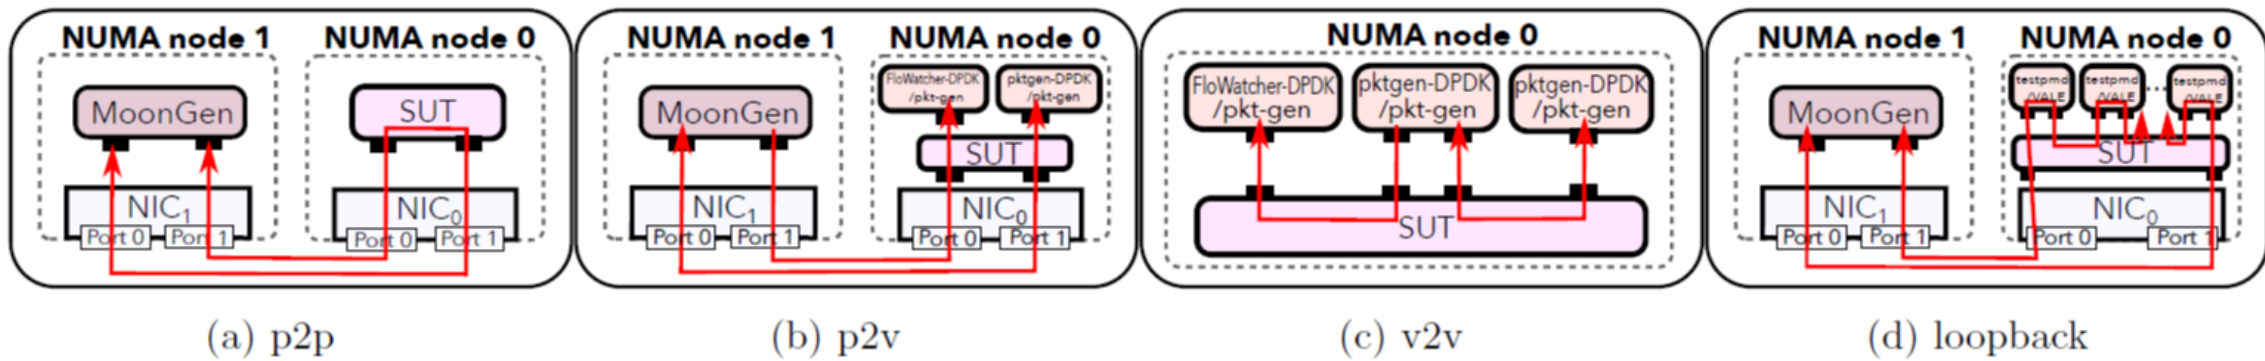
Snabbswitch

Netmap/VALE

# Tools

- Traffic generators and monitors @ 10Gbps
  - MoonGen
  - pktgen-DPDK
  - pkt-gen
  - FloWatcher-DPDK
- Hypervisors
  - QEMU/KVM + CentOS 7
  - Docker
- Virtual network functions
  - Traffic monitors
  - testpmd/VALE



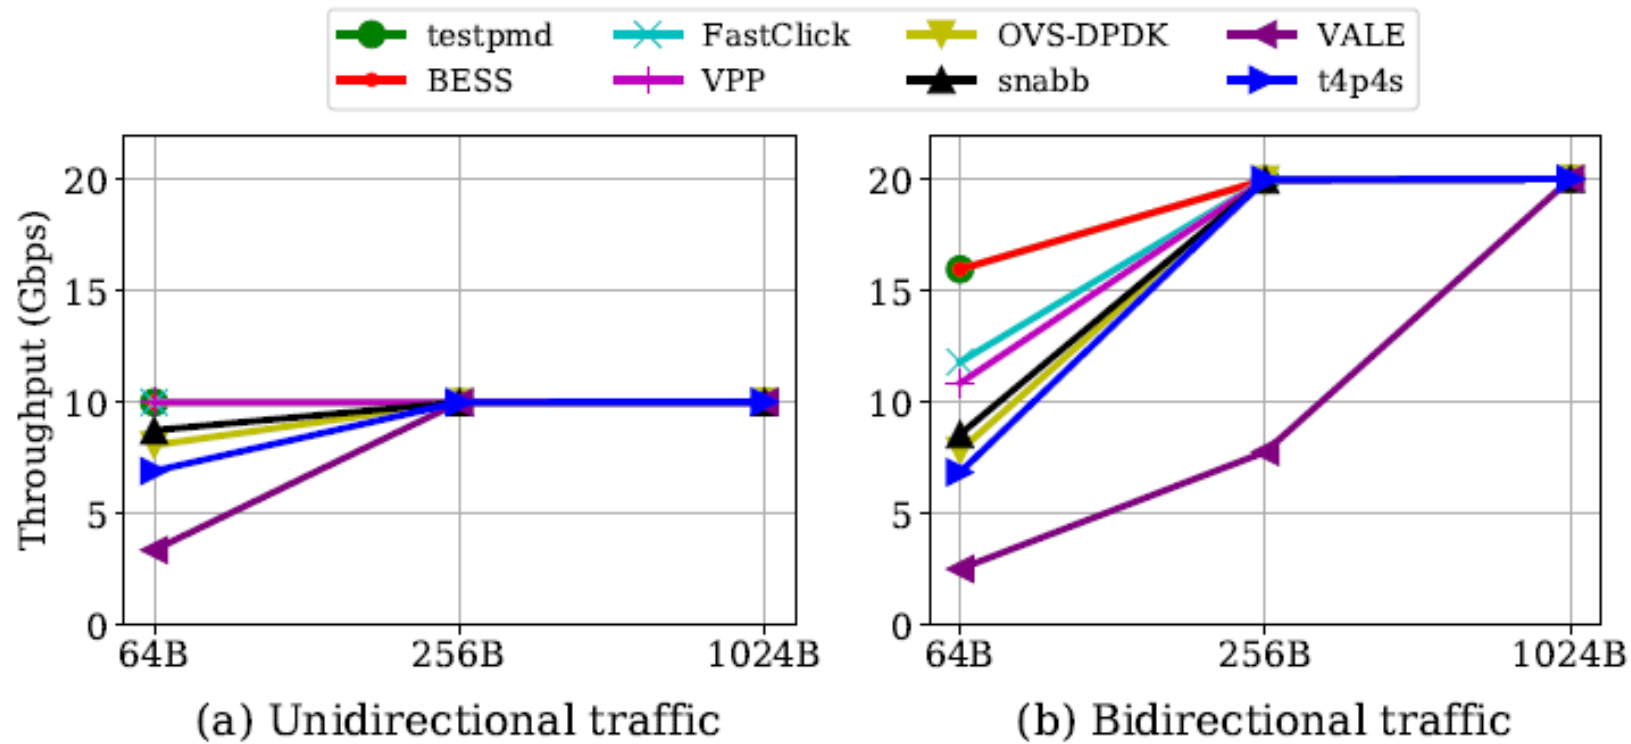


# Testbed settings

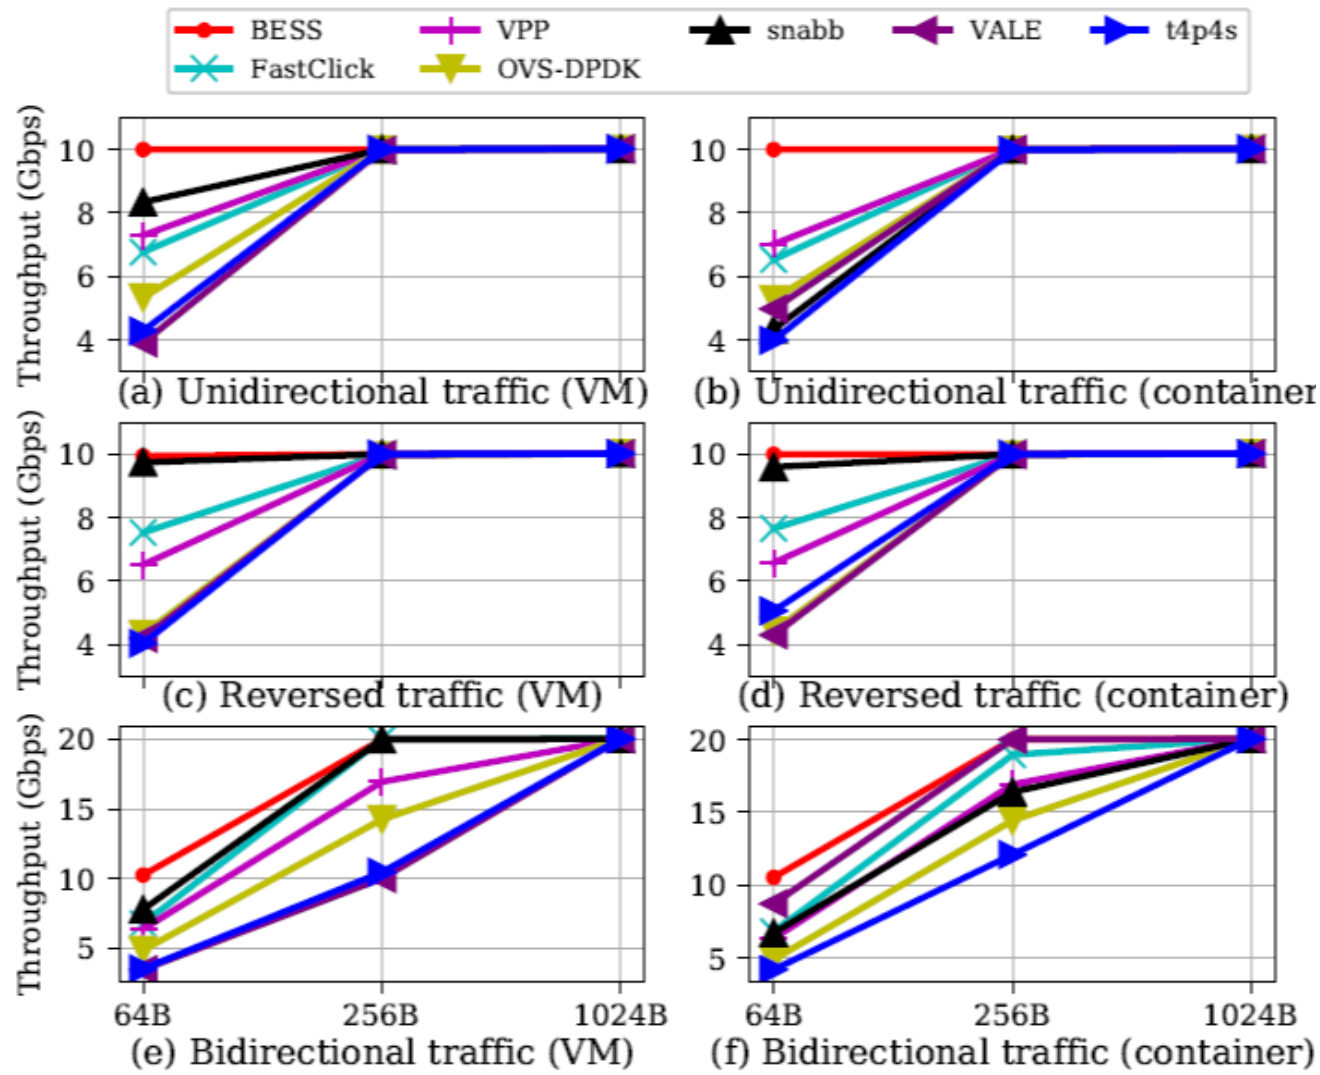
# Experimental parameters

- Traffic patterns
  - Synthetic packets with 64B, 256B, and 1024B
- CPU assignment
  - Single core for software switches
  - Four cores for each VNF
- Virtual environment
  - VM and container
- Performance metrics
  - Throughput: in Gbps, both unidirectional and bidirectional
  - Latency: Round-Trip time under 0.1/0.5/0.99 of the maximal forwarding rate

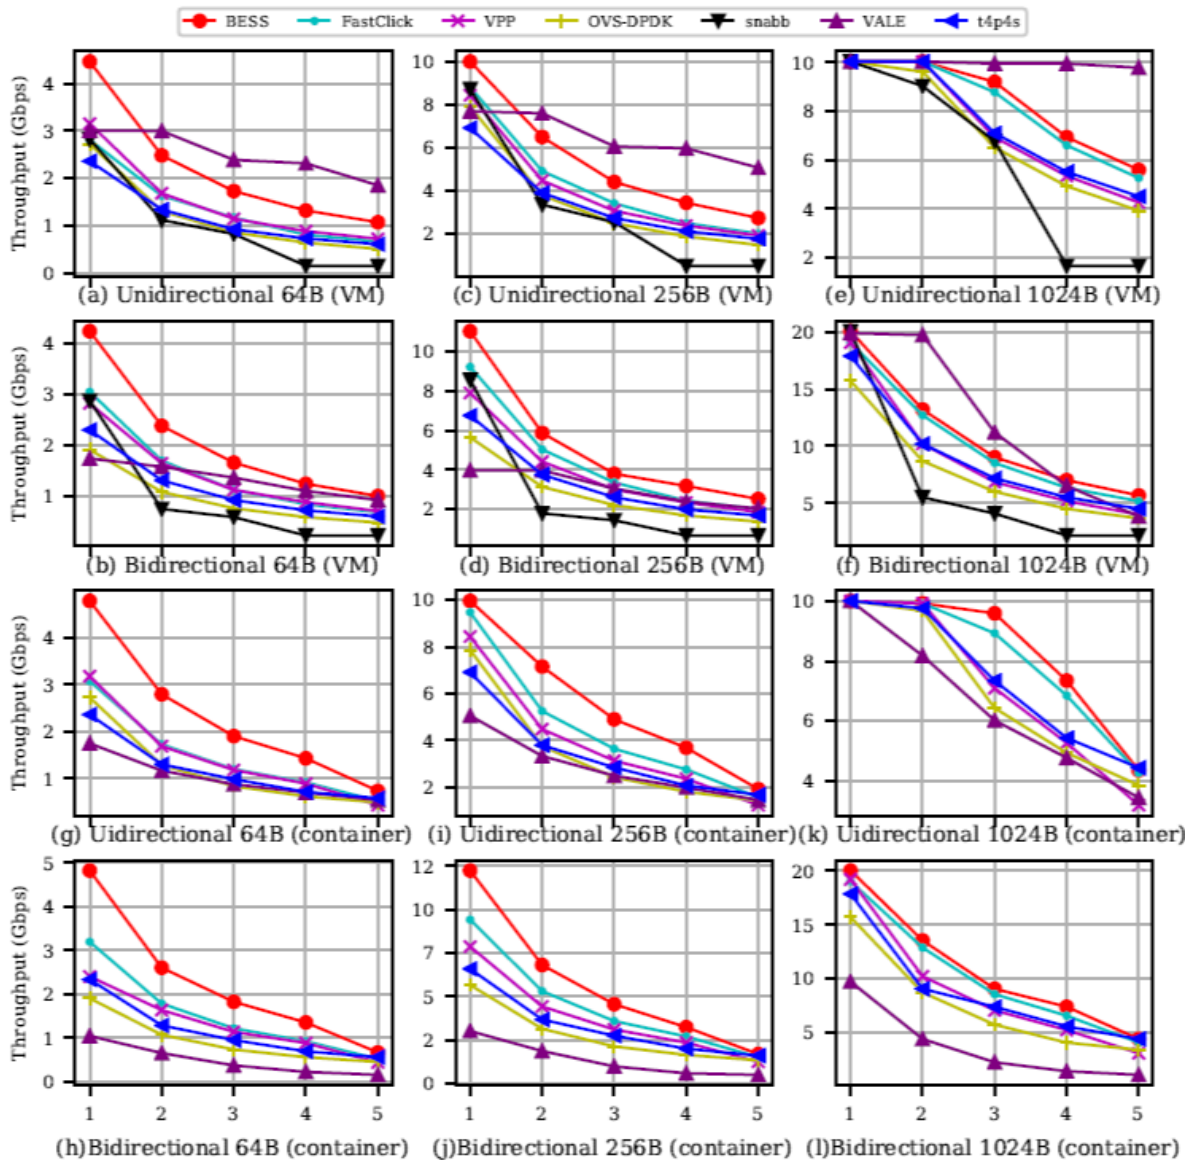
# Experimental results



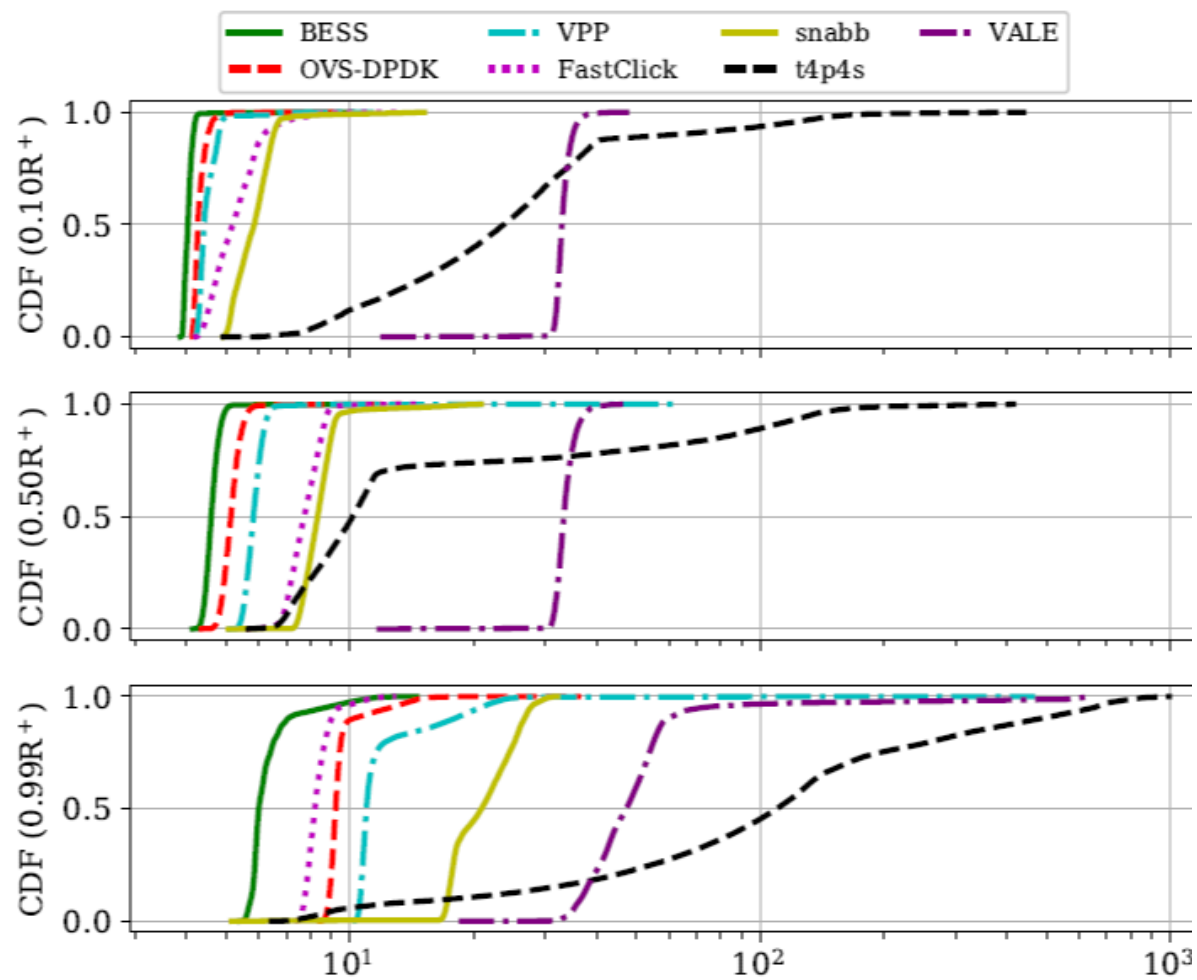
# Throughput test for p2p scenario



Throughput  
test for p2v  
scenario

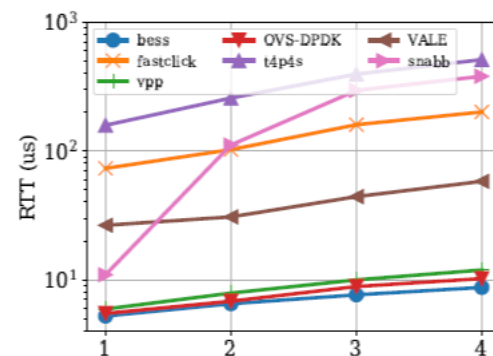


Throughput  
test for  
loopback  
scenario




Latency  
test for p2p  
scenario

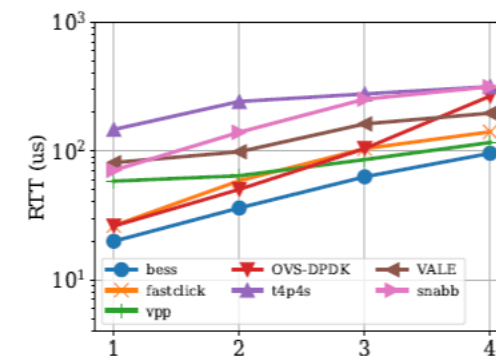
# Latency test for loopback scenario




(a)  $0.10R^+$  with VNFs deployed in VMs



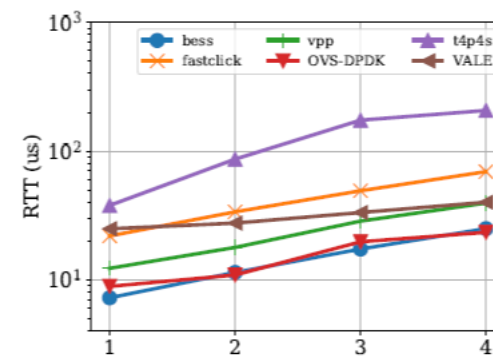
(b)  $0.50R^+$  with VNFs deployed in VMs




(c)  $0.99R^+$  with VNFs deployed in VMs



(d)  $0.10R^+$  with containerized VNFs



(e)  $0.50R^+$  with containerized VNFs



(f)  $0.99R^+$  with containerized VNFs




# Use case summary

	Best use cases	Remarks
BESS	Forwarding between physical NICs and containers	Chaining of containerized VNFs
Snabb	Fast deployment, runtime optimization	Bottlenecked with multiple VNFs
OVS-DPDK	Stateless SDN deployments	Supports OpenFlow protocol
FastClick	VNF chaining	Flexible live migration, high latency at low workload
VPP	VNF chaining	Flexible live migration
VALE	VNF chaining with high workload	Limited traffic classification and live migration capability
t4p4s	Stateful SDN deployments	Supports P4 semantics

# Conclusion

- Software switches are widely used by NFV frameworks
- It calls for a benchmarking methodology to assess their performance
- Our approach provides a comprehensive evaluation of software switches in a typical NFV environment
- Quantitative measurement with qualitative evaluation
- Code reproducible and available on GitHub





Thanks!!!

